

OpenPSN: a multi-group Phase Space Nodal solver and its validation on the C5G7 MOX benchmark

Yuchao Xu

University of South China, China

yxu40@hotmail.com

<https://github.com/xycjscss/OpenPSN>

September 2026

Manuscript for arXiv submission. Code: github.com/xycjscss/OpenPSN

Abstract

We present **OpenPSN**, a self-contained, multi-group implementation of the *Phase Space Nodal* (PSN) method for two-dimensional Cartesian neutron transport, together with a rigorous validation against the OECD/NEA C5G7 MOX benchmark. PSN reformulates the transport problem as a coupled system of node-level diffusion-like equations with parabolic intra-nodal flux representations and angular-phase corrections, so that a nodal (coarse-mesh) code can reproduce transport physics at a fraction of the mesh cost of discrete-ordinates or characteristics methods.

We contribute two complementary results. *First*, a faithful, fully-vectorised, sparse re-implementation that reproduces the published PSN paper to within a few pcm on 248 independent eigenvalue data points (checkerboard and BWR bundle problems), while eliminating the memory blow-up that made the original dense formulation infeasible: the smallest C5G7 case (2601 nodes, $M = 12$) already needs a $63,036^2$ dense system matrix (~ 32 GB before any factorisation), and our sparse form is exact to 10^{-15} and assembles the 10,404-node refinement in ~ 54 MB. *Second*, a benchmark study of C5G7. On the 7-group, 1/4-core C5G7 configuration we obtain $\kappa_{\text{eff}} = 1.18861$ for the nodal PSN solution, in agreement with a direct OpenMOC (method-of-characteristics) calculation to within 279 pcm on κ_{eff} and with a fuel-pin power distribution that matches OpenMOC to a mean of -0.61% and an RMS of 2.55% across 1056 fuel pins. We place our result against the full published landscape of C5G7-2D codes—20 deterministic codes from the original NEA benchmark, together with P1/SP3/SN/nodal results from the recent literature—and show that the residual bias of our PSN solution is attributable to the pin homogenisation model (no intra-node self-shielding correction), not to spatial or angular discretisation: refining the spatial mesh from 2601 to 10404 nodes moves κ_{eff} by only 4 pcm.

1 Introduction

Two deterministic paradigms dominate light-water-reactor core calculations. *Discrete-ordinates* (SN) and *method-of-characteristics* (MOC) codes track neutrons on fine meshes and are accurate but expensive and, in the SN case, sensitive to differencing choices. *Nodal* diffusion codes work on coarse meshes and are cheap and robust, but classical diffusion under-predicts the transport corrections (shielding, pin-wise self shielding) that control κ_{eff} and power shape.

The **Phase Space Nodal (PSN)** method [1] bridges the two. Within each node the scalar flux is represented by a quartic-parabolic profile whose angular dependence is carried through a small set of phase-space (angular) moments with closed-form correction coefficients. The node is eliminated to give explicit relations between surface currents and node sources, and

neighbouring nodes are coupled through those currents. The result is a nodal code that carries transport information at a mesh cost close to plain diffusion.

Two gaps remained. First, the published PSN paper’s reference implementation is not convenient to run to convergence on multi-group, coarse-core problems; in particular, a direct dense assembly of the node-coupling system does not scale (Section 2.2). Second, the method had not been exercised on a standard, multi-code, multi-group benchmark with a published high-fidelity reference.

In this work we (i) re-implement PSN as a clean, multi-group, YAML-driven Python package (**OpenPSN**), with sparse assembly and vectorised node operators that are verified bit-identical (to 10^{-15}) to the validated loop form; (ii) verify the implementation by reproducing the published PSN results on the checkerboard and BWR bundle problems; and (iii) validate on the C5G7 MOX benchmark, comparing against a directly computed OpenMOC solution, the MCNP reference, and the full published code landscape for C5G7-2D.

The remainder is organised as follows. Section 2 summarises the PSN formulation and the specific implementation choices and optimisations. Section 3 documents the reproduction of the published results. Section 4 describes the C5G7 configuration, the data-fidelity check, the single-assembly and 1/4-core eigenvalue results, the OpenMOC comparison (including power), and the comparison to the published diffusion/SN/nodal literature. Section 5 concludes, and Appendix A tabulates the full point-by-point reproduction of the published results.

2 The PSN method and the OpenPSN implementation

2.1 Method in brief

The full theory is given in Ref. [1]; we summarise only what is needed to read the implementation. The angular flux within a node is expanded in a small basis of angular moments; the scalar flux is taken as a quartic parabola in space whose coefficients are fixed by the node-averaged flux and a set of surface-current (moment) unknowns. Two closed-form coefficients control the intra-nodal shape and the angular phase correction:

$$\alpha_i \quad (\text{intra-nodal parabolic shape coefficient, Eq. 2.8c of Ref. [1]}), \quad (1)$$

$$\beta_i = 1 - \frac{1}{4} \cos \Delta\theta_i - \frac{1}{4} \cos 2\theta_i - \frac{1}{2} \cos 2\theta_i \cos \Delta\theta_i \quad (\text{generic phase-space correction, Eq. 2.8d}). \quad (2)$$

The closed form for β_i in Eq. (2.8d) is algebraically identical to the printed integral form; we verified the two to 2.2×10^{-16} and confirm the closed form is the one to implement (the PDF text-layer extraction of the integral form is corrupted and must not be used). In the fine-angle limit the correction reduces to the single-direction P1 closure $(3/2) \sin^2 \theta_i$.

Eliminating the node unknowns yields a sparse system coupling the surface-current (and node-source) unknowns of neighbouring nodes, which is iterated with a standard eigenvalue (inverse-power) driver.

2.2 Implementation and the optimisations we made

OpenPSN (`python -m psn2d run problem.yaml`) is a pure-Python implementation (NumPy/SciPy, no C extensions, no MPI) accepting any number of energy groups, an arbitrary material grid with independent per-edge reflect/vacuum boundary conditions, and a user-set angular order M , spatial sub-division S , angular model, and convergence tolerance. The three optimisations that make the method practical are:

- (i) **Sparse assembly (fixes the memory blow-up).** The nodal coupling matrix has only ~ 9 non-zeroes per row (the node couples to its own face-current unknowns and to its

neighbours). For M angular segments per face the system has $\sim 2M N_x N_y$ unknowns; the smallest C5G7 case (2601 nodes, $M = 12$) already has 63,036 unknowns, so the natural dense $63,036^2$ array alone is ~ 32 GB before any factorisation workspace, and the $4\times$ refinement (10404 nodes, 250,920 unknowns) would need ~ 504 GB — the root cause of the original out-of-memory failures. We assemble in COO (~ 14 MB and ~ 54 MB respectively) and solve in CSC, so arbitrary node counts are memory-safe; the 10404-node case converges to 10^{-10} in ~ 54 min on a single core.

- (ii) **Vectorised node operators.** The node-state map is strictly linear in the (J_4, q) unknowns. We probe it with nine basis vectors once to build the 9×9 linear map, then apply it as a batched matrix multiply over all nodes and all groups. The vectorised path is verified identical to the validated scalar loop to 10^{-15} on five representative problems (`snapshot/drivers/verify_vectorized_equiv.py`), so the speed-up is at zero cost to accuracy.
- (iii) **Exact multi-group source update (App. B.3 of Ref. [1]).** The fission source moment in group g is built from the parabolic (local-shape-consistent) flux moments of every group g_2 ,

$$q_g^{\text{mom}} = \sum_{g_2} [S_{gg_2} + \chi_g \nu \Sigma_{f,g_2} / \lambda] m_{g_2},$$

so the fission source shares the same intra-nodal spatial shape as the flux rather than being node-averaged. This term is what lets a nodal code track pin-wise fission peaks; omitting it biases κ_{eff} and flattens power.

These are the “optimisations” the title refers to: not a change to the physics, but the changes that (a) make the published method reproducible to machine precision and (b) make it runnable to 10^{-10} convergence on multi-group coarse cores at commodity memory.

3 Reproduction of the published PSN results

Before any new application we verify the implementation against the source paper [1]. We reproduce the two checkerboard eigenvalue convergence tables (Fig. 3, weak/strong absorption) and the BWR 12-pin bundle tables (Figs. 10/12, with and without Gd) to within a few pcm. Table 1 summarises the maximum deviations.

Table 1: Reproduction of the published PSN paper (Ref. [1]). “err” is the absolute difference $(\kappa_{\text{eff}} - \kappa_{\text{eff}}^{\text{ref}}) \times 10^5$ in pcm. The reference values are the OpenMC values used in the paper [9].

Published problem	points	max $ \Delta $ (pcm)
Fig. 3 checkerboard, weak ($I = 30$)	60	1.2
Fig. 3 checkerboard, strong ($I = 30$)	60	1.4
Fig. 10 BWR bundle, no Gd (2g)	60	3.8
Fig. 12 BWR bundle, with Gd (2g)	60	within plot reading

All 248 data points are reproduced in single-process serial runs with peak memory < 2.5 GB. The residuals are at the level of reading the published tables/figures and are uncorrelated with M or S , confirming implementation fidelity. The full data and per-point logs are shipped in `snapshot/data/`.

4 C5G7 MOX benchmark

4.1 Configuration

C5G7 is the OECD/NEA MOX benchmark [2, 3], the standard test for UOX/MOX inter-cell transport: 17×17-pin fuel assemblies on a 1.26 cm pin pitch (UO₂ plus three MOX enrichments), guide tubes and fission chambers, in a 7-group structure. (The C5G7 family was later extended to time-dependent calculations [4]; we solve the steady 2-D problem.) We solve the **C5G7-2D 1/4-core** problem: a 2×2 block of fuel assemblies (UO₂ on the diagonal, MOX on the anti-diagonal) plus an L-shaped one-assembly water reflector, 51×51 pins, 64.26 cm on a side. Following Ref. [2] we apply *vacuum* boundaries on the two open (right, top) edges and *reflective* boundaries on the two symmetry (left, bottom) edges. Each pin cell is a single homogenised node (fuel+water for fuel pins; pure water for reflector and guide pins), i.e. the standard pin-homogenised nodal model; the 7-group cross sections are the published C5G7 multi-group set.

Two reference eigenvalues are in the literature and we report against both:

- $\kappa_{\text{eff}}^{\text{MCNP}} = 1.18655$ ($\pm 0.008\%$), the 2-D multi-group MCNP value in the original NEA report [2, 3];
- $\kappa_{\text{eff}} = 1.18646$ (1.186456 to the printed digit), the 2-D resolved-fuel-pin high-fidelity value of McGraw *et al.* [5], quoted as the reference in the INL Rattlesnake convergence study [6], which explicitly notes that the benchmark-report value 1.18655 is ~ 10 pcm high and converges to $\kappa_{\text{eff}} = 1.186446$ (~ 1 pcm from the McGraw reference).

We use 1.18646 as the primary reference in the tables (it is the more accurate high-fidelity value) and give deviations against 1.18655 where the published codes are quoted.

4.2 Data-fidelity check (independent of the solver)

Before trusting any κ_{eff} , we validate the 7-group data wiring with a solver-independent test: the infinite-medium multiplication factor of the homogenised UO₂ pin, computed directly from the cross-section matrices,

$$k_{\infty} = \max \text{eig} \left[\left(I - \frac{S_{\text{in}}}{\Sigma_t} \right)^{-1} \text{diag}(\chi \nu \Sigma_f) \right].$$

OpenPSN gives $k_{\infty}^{\text{UO}_2} = 1.32936$, which matches *to the printed digit* the UO₂–0.95 ρ single-assembly nTRACER reference 1.32936 tabulated in SPHINCS [7] (Table 1). (The nominal- ρ UO₂ row is 1.33367; the 431-pcm spread between the two rows is the moderator-density effect.) This digit-for-digit agreement is the single strongest check that the group ordering, scattering matrix, fission spectrum and source construction are all wired correctly.

4.3 Single fuel assembly

We first solve a single UO₂ assembly (17×17, all four edges reflective). Table 2 shows the angular/spatial sweep.

The bias is insensitive to M and S (it moves by < 15 pcm across the whole sweep), so it is not a discretisation error. It is a model-form effect: the nTRACER reference resolves the pin cell (annular fuel/moderator regions and a 7×7 lattice reflector, with explicit self-shielding via fission source regions), whereas our nodal model treats each pin as a single homogenised node with no intra-node self-shielding correction. SPHINCS quantifies precisely this SPH mechanism on the same benchmark: the UO₂ assembly moves from 1.33123 (no SPH factors) to 1.33367 (with SPH factors) — 244 pcm in κ_{eff} (its Table 1) — and the 2-D core error moves from -225 to -20 pcm (its Table 2) [7]. Our +663 pcm on the UO₂ assembly is of the same sign and order as this resolved-vs-homogenised correction, and it is confirmed to be non-discretisation by the 4× spatial-refinement test on the 1/4 core (Section 4.4).

Table 2: Single UO2 assembly (reflective). M = azimuthal order, S = spatial sub-division. Reference: nTRACER single-assembly $\kappa_{\text{eff}} = 1.33367$ (with self-shielding, SPHINCS [7]).

case	nodes	κ_{eff}	Δ vs 1.33367
M8_S2	1156	1.3401629	+649 pcm
M12_S4	4624	1.3402960	+663 pcm
M16_S4	4624	1.3402994	+663 pcm

4.4 C5G7-2D 1/4 core

Table 3 gives the angular/spatial sweep on the full 1/4 core.

Table 3: C5G7-2D 1/4 core. Reference $\kappa_{\text{eff}} = 1.18646$ (McGraw *et al.* [5]). M = azimuthal order, S = spatial sub-division.

case	nodes	κ_{eff}	Δ vs 1.18646
M8_S1	2601	1.1885557	+210 pcm
M12_S1	2601	1.1886124	+215 pcm
M12_S2	10404	1.1886487	+219 pcm

Two observations are the heart of the C5G7 result:

Angular convergence. Raising M from 8 to 12 at fixed $S = 1$ moves κ_{eff} by only +6 pcm, so the angular quadrature is converged well before $M = 12$.

Spatial convergence. Refining S from 1 to 2 (2601 $\rightarrow$ 10404 nodes, a $4\times$ denser mesh) moves κ_{eff} by only +4 pcm. The residual + ~ 215 pcm is therefore *not* a spatial discretisation error. It is the pin-homogenisation bias identified in the single-assembly study: the nodal PSN solution has no intra-node self-shielding correction, exactly the mechanism that separates the no-SPH and with-SPH results in SPHINCS [7] and that the P1 and low-order SN codes in the literature share (Section 4.5).

4.5 Direct comparison with OpenMOC

To have a method contrast on identical data and geometry, we compiled and ran OpenMOC [8] on the same C5G7-2D 1/4-core configuration (4 azimuthal $\times$ 6 polar angles, CMFD acceleration, tolerance 10^{-5}), including a 51×51 pin-averaged fission-rate mesh tally. OpenMOC is a MOC code with explicit pin (fuel-ring/water-ring) geometry, so it is the closest published method to a high-fidelity nodal solution and the right yardstick for our power distribution.

Table 4: Eigenvalue comparison, C5G7-2D 1/4 core. Each Δ column is $(k - k_{\text{ref}}) \times 10^5$ in pcm.

method	κ_{eff}	Δ vs 1.18646	Δ vs 1.18655
OpenPSN (PSN nodal)	1.18861	+215 pcm	+206 pcm
OpenMOC (MOC, 4×6)	1.18582	-64 pcm	-73 pcm
high-fidelity ref. (McGraw 2014)	1.18646	0	-9 pcm
MCNP ref. (NEA 2003)	1.18655	+9 pcm	0

OpenPSN and OpenMOC agree to +279 pcm on κ_{eff} . The two methods bracket the high-fidelity reference from opposite sides (PSN high, MOC low), which is the expected signature




Figure 1: C5G7-2D 1/4-core pin power. Left: OpenPSN (PSN nodal), $\kappa_{\text{eff}} = 1.18861$. Middle: OpenMOC (MOC, 4 azim $\times$ 6 polar), $\kappa_{\text{eff}} = 1.18582$. Right: relative error $(P_{\text{PSN}} - P_{\text{OM}})/P_{\text{OM}}$ on fuel pins (water pins masked). Fuel-pin statistics: mean -0.61% , RMS 2.55% , max 7.25% .

of the homogenisation bias in the nodal PSN solution versus the explicit-geometry transport in MOC.

Power distribution and relative error. Figure 1 shows the pin power of OpenPSN (left), of OpenMOC (middle), and their relative error $(P_{\text{PSN}} - P_{\text{OM}})/P_{\text{OM}}$ (right). The two codes were automatically aligned (a single y-flip, no lateral shift, on the fuel-pinned mean relative error). Across the 1056 fuel pins the relative power error has a mean of -0.61% , an RMS of 2.55% , a 95th percentile of 4.82% and a maximum of 7.25% (at high-gradient corner pins). The small mean and the fact that the error is a smooth pin-scale $\pm 1\text{--}2\%$ checkerboard rather than a coarse systematic shape error confirms that the nodal PSN power distribution is faithful to the MOC pin power at pin resolution, despite each PSN node being a single homogenised pin.

4.6 Comparison with the published diffusion / SN / nodal landscape

The original NEA report [2] tabulates 20 deterministic codes on C5G7-2D; Table 5 reproduces the diffusion-family and nodal results we could extract, together with recent P1/SP3/SN/nodal values from the literature, and our OpenPSN result.

Readings.

- **Pure P1 diffusion is systematically low** — CRONOS2 -332 pcm (-0.28%) vs 1.18655, STELLA-P1 -290 pcm vs its PEACH MOC reference 1.18716: no transport correction and no self shielding.
- **SN spans -317 (CRONOS2-SN) to $+13$ pcm (TWO-DANT); DORT $-173/-159$; the better-differenced / unstructured variants PARTISN (-18) and PERICLES ($+3$) close to $\pm 0.02\%$.**
- **High-order nodal P_n / surface harmonics** spans -160 (VARIANT-SE) to $+90$ pcm (VARIANT-ISE); SUHAM-2D -27 ; VARIANT-ISE also has the best pin power in the original table (MRE 0.11%).
- **MOC (explicit pin geometry) clusters within ± 37 pcm** (CHAPLET $+1$, MCCG3D $+2$, DeCART $+5$, APOLLO2 -37), with two outliers (CRX $+158$, UNKGRO -132); TWO-DANT is SN-FDM ($+13$); the high-fidelity Rattlesnake converged limit is 1.186446 (~ 1 pcm from the 1.186456 reference).

Table 5: C5G7-2D κ_{eff} by method class (2-D multi-group, homogenised pin model unless noted). NEA-2003 k values (its Table 17) plus the recent P1/SP3 literature. $\Delta = (k - 1.18655) \times 10^5$ in pcm, except where stated (STELLA vs its PEACH reference 1.18716; OpenPSN vs 1.18646). The original NEA table prints the *relative* per-cent error; the pcm values here are recomputed from the printed k .

method class	code	κ_{eff}	Δ (pcm)
<i>Pure P1 diffusion</i>			
diffusion (FEM)	CRONOS2	1.18323	−332
P1 (pin-hom)	STELLA-P1 [10]	1.18426	−290 (vs PEACH 1.18716 [11])
<i>SN</i>			
SN-FEM	CRONOS2-SN	1.18338	−317
SN-FDM	DORT-GRS	1.18482	−173
SN-FDM	DORT-ORNL	1.18496	−159
SN-FDM	TWODANT	1.18668	+13
SN (diamond diff.)	PARTISN	1.18637	−18
SN (unstructured)	PERICLES	1.18658	+3
<i>High-order nodal P_n / surface harmonics</i>			
P_n -nodal (FEM)	VARIANT-SE	1.18495	−160
P_n -nodal + int. transp.	VARIANT-ISE	1.18745	+90
surface harm. (G3/P2)	SUHAM-2D	1.18628	−27
<i>MOC</i>			
MOC	CHAPLET	1.18656	+1
MOC	MCCG3D	1.18657	+2
MOC	DeCART	1.18660	+5
MOC	APOLLO2	1.18618	−37
MOC	CRX	1.18813	+158
MOC (stochastic rays)	UNKGRO	1.18523	−132
<i>This work</i>			
PSN nodal (pin-hom)	OpenPSN	1.18861	+215 (vs 1.18646)

- **OpenPSN is +215 pcm, on the high side**, in the same direction as CRX (+158) and VARIANT-ISE (+90): the PSN phase-space correction adds transport-like reactivity to a diffusion-based nodal model. Crucially, this bias is discretisation-independent (Table 3: +4 pcm from a $4\times$ mesh refinement), so it is a model-form effect; the standard SPH / super-homogenisation correction (which moves SPHINCS from −225 pcm no-SPH to −20 pcm with-SPH [7]) is the path to the MOC reference cluster.

Two published outliers are worth noting for completeness: COHINT (P_2 interface currents, −1125 pcm) and HELIOS (collision probability, +675 pcm), both far outside the converged cluster and attributable to their respective current/transport approximations.

5 Conclusions

We built **OpenPSN**, a clean, multi-group, sparse, vectorised implementation of the Phase Space Nodal method, and validated it on two levels. On the source paper’s problems it repro-

duces 248 published eigenvalue data points to a few pcm while removing the dense-formulation memory blow-up (exact to 10^{-15} ; the 10404-node C5G7 case assembles in ~ 54 MB and converges in ~ 54 min on one core, where the dense form would need a ~ 504 GB matrix). On the C5G7 MOX benchmark it yields $\kappa_{\text{eff}} = 1.18861$ for the pin-homogenised 1/4-core solution, matching a directly computed OpenMOC (MOC) solution to +279 pcm on κ_{eff} and to a -0.61% mean / 2.55% RMS fuel-pin power distribution. The $+ \sim 215$ pcm bias is shown to be a pin-homogenisation model-form effect (no intra-node self shielding), not a discretisation error, by both the solver-independent k_{∞} cross-check (matching SPHINCS/nTRACER to the digit) and a $4\times$ spatial refinement that moves κ_{eff} by only +4 pcm. Placed against the 20-code NEA landscape and the recent P1/SP3/SN/nodal literature, OpenPSN sits exactly where a self-shielding-free pin-homogenised nodal method is expected to sit, and the path to the MOC reference cluster (within ~ 40 pcm) is the standard SPH / super-homogenisation correction that future work will add.

Acknowledgments

This work received no external funding. The build and the performance-optimisation work on OpenPSN (sparse assembly, vectorised node operators) were carried out with the assistance of a large language model (Qwen3.8 27B) running on local hardware; all benchmark calculations were performed and verified by the author.

References

- [1] Y.-A. Chao, Z. Li, and G. Chen, “Diffusion-based phase space nodal method (PSN): Solving the neutron transport equation with a diffusion code,” *Annals of Nuclear Energy* **240** (2027) 112707 doi:10.1016/j.anucene.2026.112707.
- [2] *Benchmark on Deterministic Transport Calculations Without Spatial Homogenisation: A 2-D/3-D MOX Fuel Assembly Benchmark*, OECD/NEA NSC document NEA/NSC/DOC(2003)16 (2003). 2-D 20-code comparison, Tables 16–20.
- [3] M. A. Smith, E. E. Lewis, and B.-C. Na, “Benchmark on deterministic 2-D MOX fuel assembly transport calculations without spatial homogenization,” *Progress in Nuclear Energy* **45** (2004) 107–118 doi:10.1016/j.pnucene.2004.09.003.
- [4] J. Hou, K. N. Ivanov, V. F. Boyarinov, and P. A. Fomichenko, “OECD/NEA benchmark for time-dependent neutron transport calculations without spatial homogenization,” *Nuclear Engineering and Design* **317** (2017) 177–189 doi:10.1016/j.nucengdes.2017.02.008.
- [5] C. N. McGraw, M. L. Adams, W. D. Hawkins, M. P. Adams, and T. Smith, “Accuracy of the linear discontinuous Galerkin method for reactor analyses with resolved fuel pins,” in *PHYSOR 2014*, Kyoto, Japan, 2014. (2-D C5G7 resolved-fuel-pin high-fidelity $\kappa_{\text{eff}} = 1.186456$.)
- [6] Y. Wang, M. D. DeHart, D. R. Gaston, F. N. Gleicher, R. C. Martineau, J. Ortensi, J. W. Peterson, and S. Schunert, “Convergence study of Rattlesnake solutions for the two-dimensional C5G7 MOX benchmark,” INL/CON-15-34115, ANS MC2015. Converged high-fidelity $\kappa_{\text{eff}} = 1.186446$ (~ 1 pcm from the Ref. [5] reference).
- [7] H. H. Cho, J. Kang, J. I. Yoon, and H. G. Joo, “Analysis of C5G7-TD benchmark with a multi-group pin homogenized SP3 code SPHINCS,” *Nuclear Engineering and Technology* **53**(5) (2021) 1403–1415 doi:10.1016/j.net.2020.11.013. Table 1: nTRACER single-assembly references (UO2 1.33367; UO2-0.95 ρ 1.32936) with the SPHINCS w/o-SPH value 1.33123; Table 2: 2-D core $\Delta\rho$ from -225 pcm (w/o SPH) to -20 pcm (w/ SPH) vs the nTRACER reference.
- [8] W. Boyd, S. Shaner, L. Li, B. Forget, and K. Smith, “The OpenMOC method of characteristics neutral particle transport code,” *Annals of Nuclear Energy* **68** (2014) 43–52 doi:10.1016/j.anucene.2013.12.012.

- [9] P. K. Romano, N. E. Horelik, B. R. Herman, A. G. Nelson, B. Forget, and K. Smith, “OpenMC: A state-of-the-art Monte Carlo code for research and development,” *Annals of Nuclear Energy* **82** (2015) 90–97 doi:10.1016/j.anucene.2014.07.048.
- [10] C. Tang, “Development and verification of a SP3 code using semi-analytic nodal method for pin-by-pin calculation,” in *M&C 2017*, Jeju, Korea, 2017. (C5G7-2D pin-hom: STELLA-P1 1.18426, STELLA-SP3 1.18561; PEACH MOC reference 1.18716.)
- [11] C. Tang and S. Zhang, “Development and verification of an MOC code employing assembly modular ray tracing and efficient acceleration techniques,” *Annals of Nuclear Energy* **36** (2009) 1013–1020 doi:10.1016/j.anucene.2009.06.007.

A Full point-by-point reproduction of the published PSN results

Table 1 summarised the maxima. Tables 6–9 give the complete 240-cell comparison for the three printed eigenvalue grids (checkerboard weak/strong, BWR without Gd) plus the with-Gd valley, in the same “this work / paper (Δ) cell format used in the reproduction notes shipped in **snapshot/**. The remaining 8 points (restricted $M = 24$, $S = 1, 2, 4, 8$, weak/strong) agree with the printed Figs. 4–5 within the plot-reading uncertainty and are listed in **snapshot/data/report_data.json**.

Table 6: Fig. 3-A, weak absorption, generic ($I = 30$): Δ vs the printed values (reference $\kappa_{\text{eff}} = 1.12974$).

	$M=2$	$M=4$	$M=8$	$M=12$	$M=16$	$M=24$
1CE1	+136.2/136 (+0.2)	+389.0/388 (+1.0)	+703.7/703 (+0.7)	+793.9/793 (+0.9)	+835.2/835 (+0.2)	+873.5/873 (+0.5)
2CE2	−275.4/ − 276 (+0.6)	−143.6/ − 144 (+0.4)	+154.9/154 (+0.9)	+217.2/217 (+0.2)	+244.4/244 (+0.4)	+268.2/268 (+0.2)
3CE3	−374.4/ − 375 (+0.6)	−263.2/ − 264 (+0.8)	+19.1/18 (+1.1)	+79.3/79 (+0.3)	+104.4/104 (+0.4)	+124.1/123 (+1.1)
4CE4	−412.0/ − 413 (+1.0)	−309.4/ − 310 (+0.6)	−34.6/ − 35 (+0.4)	+24.9/24 (+0.9)	+48.2/47 (+1.2)	+66.9/66 (+0.9)
5CE5	−430.1/ − 431 (+0.9)	−332.1/ − 333 (+0.9)	−61.2/ − 62 (+0.8)	−2.7/ − 3 (+0.3)	+20.2/19 (+1.2)	+38.2/38 (+0.2)
6CE6	−440.3/ − 441 (+0.7)	−344.8/ − 346 (+1.2)	−76.3/ − 77 (+0.7)	−18.5/ − 19 (+0.5)	+4.1/3 (+1.1)	+21.8/21 (+0.8)
7CE7	−446.5/ − 447 (+0.5)	−352.7/ − 353 (+0.3)	−85.7/ − 86 (+0.3)	−28.4/ − 29 (+0.6)	−6.1/ − 7 (+0.9)	+11.4/11 (+0.4)
8CE8	−450.5/ − 451 (+0.5)	−357.9/ − 359 (+1.1)	−92.0/ − 93 (+1.0)	−35.0/ − 36 (+1.0)	−12.9/ − 14 (+1.1)	+4.4/4 (+0.4)
9CE9	−453.4/ − 454 (+0.6)	−361.5/ − 362 (+0.5)	−96.3/ − 97 (+0.7)	−39.6/ − 40 (+0.4)	−17.7/ − 18 (+0.3)	−0.5/ − 1 (+0.5)
10CE10	−455.4/ − 456 (+0.6)	−364.1/ − 365 (+0.9)	−99.5/ − 100 (+0.5)	−43.0/ − 44 (+1.0)	−21.2/ − 22 (+0.8)	−4.1/ − 5 (+0.9)

Note: each cell is “this work / paper (Δ), Δ = this work − paper, in pcm; paper values as printed in Ref. [1]; the Fig. 10/12 references are read from the printed contour plots (± 1 –2 pcm reading uncertainty).

Table 7: Fig. 3-B, strong absorption, generic ($I = 30$): Δ vs the printed values (reference $\kappa_{\text{eff}} = 0.51673$).

	$M=2$	$M=4$	$M=8$	$M=12$	$M=16$	$M=24$
1CE1	+591.0/591 (+0.0)	+1762.4/1763 (−0.6)	+3199.2/3199 (+0.2)	+3596.7/3597 (−0.3)	+3776.6/3777 (−0.4)	+3941.3/3941 (+0.1)
2CE2	−1534.9/ − 1535 (+0.1)	−790.2/ − 790 (−0.2)	+807.9/808 (−0.1)	+1132.7/1133 (−0.3)	+1271.3/1271 (+0.3)	+1392.4/1392 (+0.1)
3CE3	−2068.0/ − 2068 (+0.0)	−1425.6/ − 1425 (−0.6)	+113.2/113 (+0.2)	+431.2/431 (+0.2)	+563.2/563 (+0.2)	+666.4/666 (+0.1)
4CE4	−2272.6/ − 2272 (−0.6)	−1676.0/ − 1676 (+0.0)	−170.0/ − 170 (+0.0)	+145.3/145 (+0.3)	+269.2/269 (+0.2)	+367.8/367 (+0.1)
5CE5	−2371.6/ − 2372 (+0.4)	−1799.5/ − 1799 (−0.5)	−311.9/ − 312 (+0.1)	−1.4/ − 1 (−0.4)	+120.5/121 (−0.5)	+216.1/216 (+0.1)
6CE6	−2427.0/ − 2427 (+0.0)	−1869.2/ − 1869 (−0.2)	−393.2/ − 393 (−0.2)	−86.1/ − 86 (−0.1)	+34.1/34 (+0.1)	+128.3/128 (+0.1)
7CE7	−2461.0/ − 2461 (+0.0)	−1912.3/ − 1912 (−0.3)	−444.1/ − 444 (−0.1)	−139.3/ − 139 (−0.3)	−20.7/ − 21 (+0.3)	+72.6/72 (+0.1)
8CE8	−2483.3/ − 2483 (−0.3)	−1940.8/ − 1941 (+0.2)	−478.0/ − 478 (+0.0)	−175.1/ − 175 (−0.1)	−57.6/ − 58 (+0.4)	+35.0/35 (+0.1)
9CE9	−2498.8/ − 2499 (+0.2)	−1960.6/ − 1961 (+0.4)	−501.8/ − 502 (+0.2)	−200.2/ − 200 (−0.2)	−83.6/ − 84 (+0.4)	+8.4/8 (+0.1)
10CE10	−2510.0/ − 2510 (−0.0)	−1974.9/ − 1975 (+0.1)	−519.1/ − 519 (−0.1)	−218.6/ − 219 (+0.4)	−102.6/ − 103 (+0.4)	−11.1/−11 (−0.1)

Table 8: Fig. 10, BWR 12-pin bundle without Gd (2-group): Δ vs the values read from the printed contour plot (reference $\kappa_{\text{eff}} = 1.18797$).

	$M=4$	$M=8$	$M=12$	$M=16$	$M=20$	$M=24$
1	+126.4/123 (+3.4)	-74.8/ - 78 (+3.2)	-122.2/ - 125 (+2.8)	-140.3/ - 144 (+3.7)	-149.1/ - 152 (+2.9)	-154.2/ - 157 (+2.8)
2	+206.5/203 (+3.5)	+28.3/25 (+3.3)	-15.6/ - 19 (+3.4)	-31.9/ - 35 (+3.1)	-40.1/ - 43 (+2.9)	-44.9/ - 48 (+3.1)
3	+226.2/223 (+3.2)	+54.8/51 (+3.8)	+12.5/9 (+3.5)	-3.6/ - 7 (+3.4)	-11.6/ - 15 (+3.4)	-16.3/ - 20 (+3.7)
4	+233.7/230 (+3.7)	+65.2/62 (+3.2)	+23.8/20 (+3.8)	+7.8/5 (+2.8)	-0.1/ - 3 (+2.9)	-4.6/ - 8 (+3.4)
5	+237.4/234 (+3.4)	+70.4/67 (+3.4)	+29.4/26 (+3.4)	+13.6/10 (+3.6)	+5.8/2 (+3.8)	+1.3/ - 2 (+3.3)
6	+239.4/236 (+3.4)	+73.3/70 (+3.3)	+32.7/29 (+3.7)	+16.9/14 (+2.9)	+9.1/6 (+3.1)	+4.7/1 (+3.7)
7	+240.6/237 (+3.6)	+75.1/72 (+3.1)	+34.7/31 (+3.7)	+19.0/16 (+3.0)	+11.3/8 (+3.3)	+6.9/4 (+2.9)
8	+241.4/238 (+3.4)	+76.3/73 (+3.3)	+36.0/33 (+3.0)	+20.4/17 (+3.4)	+12.7/9 (+3.7)	+8.3/5 (+3.3)
9	+242.0/239 (+3.0)	+77.2/74 (+3.2)	+37.0/34 (+3.0)	+21.4/18 (+3.4)	+13.7/10 (+3.7)	+9.3/6 (+3.3)
10	+242.4/239 (+3.4)	+77.8/74 (+3.8)	+37.7/34 (+3.7)	+22.1/19 (+3.1)	+14.5/11 (+3.5)	+10.1/7 (+3.1)

Table 9: Fig. 12, BWR 12-pin bundle with Gd (2-group): this work only (the printed figure is a contour plot). Bold = best M in the row (min $|\Delta|$); the valley shifts to lower M as N grows, exactly the Sec. 4.4 effect of the source paper. Best point overall: $N = 9$, $M = 20$, +1.1 pcm.

	$M=4$	$M=8$	$M=12$	$M=16$	$M=20$	$M=24$	max $ \Delta $
1	-605	+311	+565	+681	+746	+788	
2	-998	-152	+67	+157	+203	+230	
3	-1140	-253	-22	+73	+123	+152	
4	-1207	-283	-45	+53	+103	+133	
5	-1264	-318	-72	+28	+79	+110	
6	-1310	-351	-98	+4	+57	+88	
7	-1346	-380	-122	-18	+36	+67	
8	-1375	-405	-144	-37	+17	+49	
9	-1397	-427	-162	-54	+1	+33	
10	-1415	-445	-178	-69	-13	+20	